

```
"cells": [
  {
    "cell_type": "markdown",
    "metadata": {},
    "source": [
      "# Veterans First AI: Building a RAG-Enabled Conversational
      Model\n",
      "\n",
      "This notebook provides a complete, end-to-end implementation for
      building the \"Veterans First\" AI assistant, as detailed in the
      comprehensive step-by-step guide. We will build a
      **Retrieval-Augmented Generation (RAG)** system, which grounds a
      fine-tuned Large Language Model (LLM) in a custom knowledge base of
      VA-related documents. This ensures the AI's answers are accurate,
      relevant, and drawn from authoritative sources.\n",
      "\n",
      "#### Project Architecture\n",
      "\n",
      "The system follows a RAG pipeline to provide accurate,
      context-aware answers:\n",
      "\n",
      "1. **User Query:** A user asks a question through the web
      interface.\n",
      "2. **Backend Server:** A Flask server receives the query.\n",
      "3. **Semantic Search:** The server converts the query into an
      embedding and searches a FAISS vector database (our knowledge base) to
      find the most relevant text chunks from VA laws, manuals, and
      decisions.\n",
      "4. **Prompt Augmentation:** The retrieved text is combined with
      the user's original query into a detailed prompt.\n",
      "5. **LLM Call:** This augmented prompt is sent to our fine-tuned
      OpenAI model. The fine-tuning helps the model understand the specific
      domain language, tone, and desired response format.\n",
      "6. **Response:** The model generates a response grounded in the
      provided context, which is then sent back to the user.\n",
      "\n",
      ""
    ]
  },
  {
    "cell_type": "markdown",
    "metadata": {},
    "source": [
      "## Step 0: Setup and Installation\n",
      "\n",
      "First, we'll install all the necessary Python libraries for this
      project and set up our environment variables. You will need an OpenAI
      API key for this notebook to work."
    ]
  }
]
```

```
        ],
    },
    {
        "cell_type": "code",
        "execution_count": null,
        "metadata": {},
        "outputs": [],
        "source": [
            "# Install required libraries\n",
            "!pip install openai==1.3.3 beautifulsoup4==4.12.2\n",
            "faiss-cpu==1.7.4 sentence-transformers==2.2.2 flask==3.0.0\n",
            "flask-cors==4.0.0 pyngrok==7.0.0"
        ]
    },
    {
        "cell_type": "code",
        "execution_count": null,
        "metadata": {},
        "outputs": [],
        "source": [
            "import os\n",
            "from google.colab import userdata\n",
            "\n",
            "# Securely get the OpenAI API key from Colab's secret manager\n",
            "try:\n",
            "    os.environ['OPENAI_API_KEY'] =\n",
            "userdata.get('OPENAI_API_KEY')\n",
            "    print(\"OpenAI API key loaded successfully.\")\n",
            "except userdata.SecretNotFoundError:\n",
            "    print(\"ERROR: OpenAI API key not found. Please add it to\n",
            "your Colab secrets.\")\n",
            "    print(\"Go to the '🔑' icon on the left panel and add a new\n",
            "secret with the name 'OPENAI_API_KEY'.\")"
        ]
    },
    {
        "cell_type": "markdown",
        "metadata": {},
        "source": [
            "## Step 1: Data Collection\n",
            "\n",
            "The foundation of our AI is a comprehensive knowledge base. In\n",
            "this step, we gather the raw data, including laws, regulations, and\n",
            "procedure manuals. For this demonstration, we'll create a few sample\n",
            "text files to simulate this process. A real-world implementation would\n",
            "involve extensive scraping and downloading from sources like the eCFR,\n",
            "VA websites, and BVA decision repositories."
        ]
    }
]
```

```
},
{
  "cell_type": "code",
  "execution_count": null,
  "metadata": { },
  "outputs": [ ],
  "source": [
    "import os\n",
    "\n",
    "# Create directories for our project structure\n",
    "os.makedirs('raw_data', exist_ok=True)\n",
    "os.makedirs('processed_data', exist_ok=True)\n",
    "os.makedirs('knowledge_base', exist_ok=True)\n",
    "os.makedirs('application', exist_ok=True)\n",
    "\n",
    "# --- Create Sample Data Files ---\n",
    "\n",
    "sample_cfr_text = \"\"\"\n",
    "Title 38 CFR § 3.303 Principles of service connection.\n",
    "Service connection will be granted if the evidence demonstrates\nthat a particular injury or disease resulting in disability was\nincurred coincident with service in the Armed Forces, or if\npre-existing, was aggravated by service. Service connection connotes\nmany factors but the main elements are the fact of disease or injury\nin service and a nexus between the in-service disease or injury and\nthe present disability. The veteran's entire service medical record\nmust be reviewed.\n",
    \"\"\"\n",
    "with open('raw_data/sample_38cfr.txt', 'w') as f:\n",
    "    f.write(sample_cfr_text)\n",
    "\n",
    "sample_m21_text = \"\"\"\n",
    "M21-1 Adjudication Procedures Manual, Part V, Subpart ii, Chapter\n1.\n",
    "The Appeals Modernization Act (AMA) provides veterans with three\noptions to seek review of a VA decision: (1) Higher-Level Review\n(HLR), (2) a Supplemental Claim, or (3) an appeal to the Board of\nVeterans' Appeals (BVA). For a Supplemental Claim, the claimant must\nsubmit new and relevant evidence. A Higher-Level Review involves a de\nnovo review of the issue(s) based on the evidence of record at the\ntime of the prior decision.\n",
    \"\"\"\n",
    "with open('raw_data/sample_m21_manual.txt', 'w') as f:\n",
    "    f.write(sample_m21_text)\n",
    "\n",
    "sample_bva_decision = \"\"\"\n",
    "BVA Decision Docket No. 20-01234\n",
    "The veteran seeks service connection for sleep apnea. The service
```

treatment records are silent for complaints or treatment for sleep-disordered breathing. A post-service diagnosis shows moderate obstructive sleep apnea. However, the private medical opinion fails to provide a sufficient nexus linking the current condition to the veteran's active duty service. Therefore, the Board finds that service connection for sleep apnea is denied.\n",

```
"\"\"\n",
"with open('raw_data/sample_bva_decision.txt', 'w') as f:\n",
"    f.write(sample_bva_decision)\n",
"\n",
"print(\"Sample data files created in 'raw_data/' directory.\")"
]
},
{
"cell_type": "markdown",
"metadata": {},
"source": [
"## Step 2: Data Preparation\n",
"\n",
"Now we process the raw data. This involves two key tasks:\n",
"1. **Cleaning and Chunking:** We clean the raw text and split it into small, logical chunks. These chunks will form the documents in our searchable knowledge base.\n",
"2. **Generating Fine-Tuning Data:** We create high-quality Question-Answer pairs from the text chunks. This data will be used to teach our base model the specific style, tone, and format for answering questions about VA benefits."
]
},
{
"cell_type": "code",
"execution_count": null,
"metadata": {},
"outputs": [],
"source": [
"import json\n",
"import re\n",
"import glob\n",
"\n",
"# Configuration for this step\n",
"RAW_DATA_DIR = 'raw_data'\n",
"PROCESSED_DATA_DIR = 'processed_data'\n",
"KNOWLEDGE_BASE_DIR = os.path.join(PROCESSED_DATA_DIR,
'knowledge_chunks')\n",
"FINETUNE_OUTPUT_FILE = os.path.join(PROCESSED_DATA_DIR,
'finetune_data.jsonl')\n",
"\n",
"os.makedirs(KNOWLEDGE_BASE_DIR, exist_ok=True)\n",
```

```

"\n",
"def clean_text(text):\n",
"    """Cleans raw text by removing excessive
whitespace.\n",
"    text = re.sub(r'\s+', ' ', text).strip()\n",
"    return text\n",
"\n",
"def chunk_text(text, chunk_size=150, overlap=30):\n",
"    """Splits text into smaller, overlapping chunks.\n",
"    words = text.split()\n",
"    if not words: return []\n",
"    chunks = []\n",
"    for i in range(0, len(words), chunk_size - overlap):\n",
"        chunk = ' '.join(words[i:i + chunk_size])\n",
"        chunks.append(chunk)\n",
"    return chunks\n",
"\n",
"def generate_qa_from_chunk(chunk, source):\n",
"    """Generates a plausible Q&A pair from a text
chunk.\n",
"    first_sentence = chunk.split('.')[0].strip()\n",
"    if len(first_sentence) < 20 or len(first_sentence) > 200:
return None\n",
"\n",
"    if '\"§\"' in first_sentence:\n",
"        question = f"What does the regulation '{first_sentence}'\nstate?\n",
"        elif '\"denied\"' in chunk.lower():\n",
"            question = f"What is a common reason for denial\nmentioned in VA decisions?\n",
"        else:\n",
"            question = f"Can you explain the VA process regarding\n'{first_sentence}'?\n",
"\n",
"    answer = chunk\n",
"    return {\n",
"        \"messages\": [\n",
"            {\"role\": \"user\", \"content\": question},\n",
"            {\"role\": \"assistant\", \"content\": f"According\n"
to {source}, {answer}}\n",
"        ]\n",
"    }\n",
"\n",
"# --- Main Execution ---\n",
"finetune_examples = []\n",
"raw_files = glob.glob(os.path.join(RAW_DATA_DIR, '*.txt'))\n",
"\n",
"for filepath in raw_files:\n"

```

```

    "     filename = os.path.basename(filepath)\n",
    "     print(f\"Processing {filename}...\")\n",
    "     with open(filepath, 'r', encoding='utf-8') as f:\n",
    "         raw_text = f.read()\n",
"\n",
"         cleaned_text = clean_text(raw_text)\n",
"         text_chunks = chunk_text(cleaned_text)\n",
"\n",
"         for i, chunk in enumerate(text_chunks):\n",
"             chunk_filename =
f\"{os.path.splitext(filename)[0]}_chunk_{i+1}.txt\"\n",
"             with open(os.path.join(KNOWLEDGE_BASE_DIR,
chunk_filename), 'w', encoding='utf-8') as cf:\n",
"                 cf.write(chunk)\n",
"\n",
"             for chunk in text_chunks:\n",
"                 qa_pair = generate_qa_from_chunk(chunk, filename)\n",
"                 if qa_pair:\n",
"                     finetune_examples.append(qa_pair)\n",
"\n",
"with open(FINETUNE_OUTPUT_FILE, 'w', encoding='utf-8') as f:\n",
"    for example in finetune_examples:\n",
"        f.write(json.dumps(example) + '\\n')\n",
"\n",
"print(f\"\\nSuccessfully created {len(finetune_examples)}\n
fine-tuning examples at: {FINETUNE_OUTPUT_FILE}\")\n",
"print(f\"Knowledge base chunks saved in: {KNOWLEDGE_BASE_DIR}\")"
]
},
{
"cell_type": "markdown",
"metadata": {},
"source": [
"## Step 3: Build the Knowledge Base with Embeddings\n",
"\n",
"This is the core of our RAG system. We convert each text chunk from the previous step into a numerical representation (an \"embedding\") using an OpenAI model. These embeddings capture the semantic meaning of the text. We then store them in a FAISS vector database, which allows for incredibly fast and efficient similarity searching."
]
},
{
"cell_type": "code",
"execution_count": null,
"metadata": {},
"outputs": []
,
```

```

"source": [
    "import faiss\n",
    "import numpy as np\n",
    "import pickle\n",
    "from openai import OpenAI\n",
    "\n",
    "# --- Configuration ---\n",
    "KNOWLEDGE_CHUNKS_DIR = 'processed_data/knowledge_chunks'\n",
    "VECTOR_STORE_DIR = 'knowledge_base'\n",
    "FAISS_INDEX_FILE = os.path.join(VECTOR_STORE_DIR,
'vector_store.index')\n",
    "METADATA_FILE = os.path.join(VECTOR_STORE_DIR,
'metadata.pkl')\n",
    "\n",
    "client = OpenAI()\n",
    "\n",
    "def get_embedding(text, model=\"text-embedding-3-small\"):\n",
        """Generates an embedding for a given text.\n""",
        text = text.replace("\\\\n\\", " \")\n",
        return client.embeddings.create(input=[text],
model=model).data[0].embedding\n",
    "\n",
    "# --- Main Execution ---\n",
    "chunk_files = glob.glob(os.path.join(KNOWLEDGE_CHUNKS_DIR,
\"*.txt\"))\n",
    "embeddings = []\n",
    "metadata = []\n",
    "\n",
    "print(f\"Creating embeddings for {len(chunk_files)} text
chunks...\")\n",
    "for filepath in chunk_files:\n",
        with open(filepath, 'r', encoding='utf-8') as f:\n",
            text_content = f.read()\n",
            if text_content.strip():\n",
                embedding_vector = get_embedding(text_content)\n",
                embeddings.append(embedding_vector)\n",
                metadata.append({'content': text_content, 'source':
os.path.basename(filepath)})\n",
    "\n",
    "embedding_matrix = np.array(embeddings).astype('float32')\n",
    "d = embedding_matrix.shape[1] # Dimension of embeddings\n",
    "index = faiss.IndexFlatL2(d)\n",
    "index.add(embedding_matrix)\n",
    "\n",
    "print(f\"Saving FAISS index to {FAISS_INDEX_FILE}...\")\n",
    "faiss.write_index(index, FAISS_INDEX_FILE)\n",
    "\n",
    "print(f\"Saving metadata to {METADATA_FILE}...\")\n",

```

```

    "with open(METADATA_FILE, 'wb') as f:\n",
    "    pickle.dump(metadata, f)\n",
    "\n",
    "print(f\"\\nKnowledge base created with {index.ntotal}\n"
documents.\") "
]
},
{
"cell_type": "markdown",
"metadata": {},
"source": [
"## Step 4: Fine-Tune the Model\n",
"\n",
"With our knowledge base built, we now fine-tune a base model\n(like `gpt-3.5-turbo`) using the Q&A pairs we generated. This doesn't\nteach the model new facts—the RAG system handles that. Instead, it\nteaches the model the *style* of a helpful VA claims expert: how to\nstructure answers, what tone to use, and how to refer to sources. This\nmakes the model a better \"reasoning engine\" for our specific\ntask.\n",
"\n",
"**This step will create a fine-tuning job on OpenAI's servers. It\ncan take some time to complete (from minutes to hours depending on\ndata size). You will need to copy the final model ID for the next\nstep.**"
]
},
{
"cell_type": "code",
"execution_count": null,
"metadata": {},
"outputs": [],
"source": [
"import time\n",
"\n",
"# --- Configuration ---\n",
"FINETUNE_DATA_FILEPATH = 'processed_data/finetune_data.jsonl'\n",
"MODEL_SUFFIX = 'veterans-first-ai-colab'\n",
"\n",
"# --- Main Execution ---\n",
"print(f\"Uploading training file: {FINETUNE_DATA_FILEPATH}\")\n",
"with open(FINETUNE_DATA_FILEPATH, \"rb\") as f:\n",
"    training_file = client.files.create(file=f,\npurpose=\"fine-tune\")\n",
"print(f\"File uploaded successfully. File ID:\n{training_file.id}\")\n",
"\n",
"print(f\"\\nCreating fine-tuning job...\")\n",
]
}

```

```

"job = client.fine_tuning.jobs.create(\n",
"    training_file=training_file.id,\n",
"    model=\"gpt-3.5-turbo\", \n",
"    suffix=MODEL_SUFFIX,\n",
")\n",
"print(f\"Fine-tuning job created successfully. Job ID:\n{job.id}\")\n",
"print(\"You can monitor the job's progress on the OpenAI\nwebsite.\")\n",
"\n",
"# --- Optional: Monitor job progress in Colab ---\n",
"print(\"\\nMonitoring job status (press Ctrl+C to stop)...\")\n",
"try:\n",
"    while True:\n",
"        job_status = client.fine_tuning.jobs.retrieve(job.id)\n",
"        status = job_status.status\n",
"        print(f\"Current status: {status}\")\n",
"        if status == \"succeeded\": \n",
"            fine_tuned_model_id = job_status.fine_tuned_model\n",
"            print(f\"\\n✅ Fine-tuning succeeded! Your model ID\nis: {fine_tuned_model_id}\")\n",
"            print(\"➡ PLEASE COPY THIS ID. You will need it for\nthe next step.\")\n",
"            break\n",
"        elif status in [\"failed\", \"cancelled\"]:\n",
"            print(f\"\\n❌ Job {status}. Check the OpenAI\ndashboard for details.\")\n",
"            break\n",
"        time.sleep(60)\n",
"except KeyboardInterrupt:\n",
"    print(\"\\nMonitoring stopped. You can check the job status\non the OpenAI website.\")"
]
},
{
"cell_type": "markdown",
"metadata": {},
"source": [
"## Step 5 & 6: Run the Backend and Web Interface\n",
"\n",
"Now we'll combine everything into a functional application. The\ncode below does the following:\n",
"1. **Defines a Flask Backend:** This is a lightweight web server\nthat will load our FAISS knowledge base and handle the full RAG\npipeline.\n",
"2. **Creates a Public URL:** We use `pyngrok` to create a\nsecure, public URL for our Flask app, allowing our web interface to\naccess it from within Colab.\n",

```

```

"3. **Launches the Web UI:** We embed the HTML and JavaScript for
our chat interface directly into the notebook. The JavaScript is
configured to communicate with the `ngrok` URL of our backend.\n",
"\n",
"**IMPORTANT:** Paste the fine-tuned model ID you copied from the
previous step into the `FINE_TUNED_MODEL_ID` variable in the code cell
below."
]
},
{
"cell_type": "code",
"execution_count": null,
"metadata": {},
"outputs": [],
"source": [
"from flask import Flask, request, jsonify\n",
"from flask_cors import CORS\n",
"from pyngrok import ngrok\n",
"from IPython.display import display, HTML\n",
"import threading\n",
"\n",
"# --- PASTE YOUR FINE-TUNED MODEL ID HERE ---\n",
"FINE_TUNED_MODEL_ID = \"ft:gpt-3.5-turbo...\" # <-- REPLACE
THIS\n",
"\n",
"# --- Flask Backend Application ---\n",
"app = Flask(__name__)\n",
"CORS(app)\n",
"\n",
"# Load knowledge base\n",
"rag_index = faiss.read_index(FAISS_INDEX_FILE)\n",
"with open(METADATA_FILE, 'rb') as f:\n",
"    rag_metadata = pickle.load(f)\n",
"\n",
"@app.route('/chat', methods=['POST'])\n",
"def chat():\n",
"    data = request.json\n",
"    user_query = data.get('message')\n",
"\n",
"    # 1. Retrieve context\n",
"    query_embedding =
np.array([get_embedding(user_query)]).astype('float32')\n",
"    distances, indices = rag_index.search(query_embedding,
k=3)\n",
"    retrieved_context = [rag_metadata[i]['content'] for i in
indices[0]]\n",
"    context_str = \"\\n\\n---\\n\\n\".join(retrieved_context)\n",
"\n"
]
}

```

```

"      # 2. Construct prompt\n",
"      system_prompt = {\n",
"          \"role\": \"system\",\\n",
"          \"content\": f\"You are the Veterans First AI. Answer the\nuser's question based only on the provided\ncontext.\\nCONTEXT:\\n{context_str}\\n\",\n          }\\n",
"      messages = [system_prompt, {\"role\": \"user\", \"content\":\nuser_query}]\\n",
"\\n",
"      # 3. Call fine-tuned model\n",
"      response = client.chat.completions.create(\\n",
"          model=FINE_TUNED_MODEL_ID,\\n",
"          messages=messages,\\n",
"          temperature=0.3\\n",
"      )\\n",
"      return jsonify({\"response\":\nresponse.choices[0].message.content})\\n",
"\\n",
"def run_app():\\n",
"    app.run(port=5000)\\n",
"\\n",
"# --- Launch Backend and UI ---\\n",
"if FINE_TUNED_MODEL_ID == \"ft:gpt-3.5-turbo...\":\\n",
"    print(\"ERROR: Please paste your fine-tuned model ID into the\n'FINE_TUNED_MODEL_ID' variable.\")\\n",
"else:\\n",
"    # Start flask app in a new thread\\n",
"    flask_thread = threading.Thread(target=run_app)\\n",
"    flask_thread.daemon = True\\n",
"    flask_thread.start()\\n",
"\\n",
"    # Expose the flask app with ngrok\\n",
"    public_url = ngrok.connect(5000).public_url\\n",
"    print(f\"Backend is running at: {public_url}\")\\n",
"\\n",
"    # --- HTML and JavaScript for the Frontend ---\\n",
"    html_template = f\"\\\"\\\"\n",
"        <!DOCTYPE html>\\n",
"        <html lang=\"en\">\\n",
"            <head>\\n",
"                <meta charset=\"UTF-8\">\\n",
"                <meta name=\"viewport\" content=\"width=device-width,\ninitial-scale=1.0\">\\n",
"                <title>Veterans First AI Assistant</title>\\n",
"                <script src=\"https://cdn.tailwindcss.com\"></script>\\n",
"                <style>body {{ font-family: sans-serif; }}</style>\\n",
"            </head>\\n",

```

```

    "      <body class=\"bg-slate-100\">\n",
    "          <div class=\"bg-white shadow-lg rounded-lg max-w-4xl
mx-auto my-8 flex flex-col h-[80vh]\">\n",
    "              <header class=\"bg-slate-800 text-white p-4
rounded-t-lg\"><h1 class=\"text-xl font-bold\">Veterans First AI
Assistant</h1></header>\n",
    "                  <main id=\"chat-container\" class=\"flex-1 p-6
overflow-y-auto\"></main>\n",
    "                      <footer class=\"p-4 bg-white border-t\">\n",
    "                          <form id=\"chat-form\" class=\"flex items-center
space-x-4\">\n",
    "                              <input type=\"text\" id=\"user-input\" placeholder=\"Ask a question...\" class=\"w-full px-4 py-2 border
rounded-lg\">\n",
    "                              <button type=\"submit\" id=\"send-btn\" class=\"bg-blue-600 text-white px-6 py-2 rounded-lg
font-semibold\">Send</button>\n",
    "                      </form>\n",
    "                  </footer>\n",
    "          </div>\n",
    "          <script>\n",
    "              const BACKEND_URL = '{public_url}';\n",
    "              const chatForm =
document.getElementById('chat-form');\n",
    "              const userInput =
document.getElementById('user-input');\n",
    "              const chatContainer =
document.getElementById('chat-container'),\n",
    "\n",
    "                  function appendMessage(role, content) {{\n",
    "                      const messageDiv =
document.createElement('div');\n",
    "                      const alignment = role === 'user' ? 'text-right'
: 'text-left';\n",
    "                      const bubbleColor = role === 'user' ?
'bg-blue-500 text-white' : 'bg-slate-200 text-slate-800';\n",
    "                      messageDiv.className = `my-2 ${alignment}`;\n",
    "                      messageDiv.innerHTML = `<div class=\"inline-block
p-3 rounded-lg ${bubbleColor}\">>${content}</div>`;\n",
    "                      chatContainer.appendChild(messageDiv);\n",
    "                      chatContainer.scrollTop =
chatContainer.scrollHeight;\n",
    "                  }}\n",
    "\n",
    "                  chatForm.addEventListener('submit', async (e) =>
{\n",
    "                      e.preventDefault();\n",
    "                      const userMessage = userInput.value.trim();\n",

```

```
"                if (!userMessage) return;\n",
"                appendMessage('user', userMessage);\n",
"                userInput.value = '';\n",
"\n",
"                const response = await\nfetch(`${
BACKEND_URL}/chat`, {\n\n,
"                    method: 'POST',\n",
"                    headers: {\n            'Content-Type':\n'application/json'\n        },\n\n,
"                    body: JSON.stringify({\n            message: userMessage\n        })\n    },\n\n,
"                );\n\n",
"                const data = await response.json();\n",
"                appendMessage('assistant', data.response),\n",
"            );\n",
"        </script>\n",
"    </body>\n",
"    </html>\n",
"\n",
"    # Display the frontend\n",
"    display(HTML(html_template))\n"
]
},
{
"cell_type": "markdown",
"metadata": {},
"source": [
"## Step 7: Evaluation and Maintenance\n",
"\n",
"Building the model is just the beginning. Continuous evaluation and maintenance are crucial for ensuring the AI remains accurate and helpful over time.\n",
"\n",
"### Evaluation Plan\n",
"\n",
"1. **Functional Testing:** Test a wide range of queries covering the entire claims process, from initial filing to appeals. Verify that the RAG system retrieves the correct context and the final answers are factually accurate.\n",
"2. **User Scenario Testing:** Simulate multi-turn conversations to assess the AI's ability to maintain context and guide users through complex scenarios.\n",
"3. **Accuracy Checks:** Involve Subject Matter Experts (SMEs) to review generated answers for correctness and clarity. Measure system performance, such as response time and retrieval quality.\n",
"\n",
"### Maintenance Strategy\n",
"\n"
]
```

```
"1. **Regular Data Updates:** The world of VA benefits changes. Schedule regular updates (e.g., quarterly) to the knowledge repository by re-running the data collection and knowledge base creation scripts with new laws, manuals, and BVA decisions.\n",
"2. **Monitor Performance and Feedback:** Implement a feedback mechanism (like thumbs-up/down ratings on answers) to identify areas for improvement. Use this feedback to expand the knowledge base or add new examples to the fine-tuning dataset.\n",
"3. **Ethical Compliance:** Always include a disclaimer that the AI is an informational tool and not a substitute for legal advice. Ensure user privacy is protected and that the AI's advice is never harmful."
    ]
},
],
"metadata": {
  "colab": {
    "collapsed_sections": [],
    "name": "Veterans_First_AI_Notebook.ipynb",
    "provenance": []
  },
  "kernelspec": {
    "display_name": "Python 3",
    "name": "python3"
  },
  "language_info": {
    "name": "python"
  }
},
"nbformat": 4,
"nbformat_minor": 0
}
```